

CH 19

```
In [1]: ► 1 # Import the required modules
2 import pandas as pd
3 pd.set_option('display.max_columns', None)
4 import numpy as np
5
6 # visualization
7 import matplotlib.pyplot as plt
8 import hvplot.pandas
9 import seaborn as sns
10
11 # Machine Learning
12 from sklearn.cluster import KMeans, AgglomerativeClustering, Birch
13 from sklearn.metrics import silhouette_score, calinski_harabasz_score
14 from sklearn.manifold import TSNE
15
16 # Preprocessing
17 from sklearn.preprocessing import StandardScaler
18 from sklearn.decomposition import PCA
19
20 # suppress warnings
21 import warnings
22 warnings.filterwarnings('ignore')
```

```
In [2]: 1 # Load the data into a Pandas DataFrame
2 df_market_data = pd.read_csv(
3     "Resources/crypto_market_data.csv",
4     index_col='coin_id')
5
6 # Display sample data
7 df_market_data.head(10)
```

```
Out[2]:
```

	price_change_percentage_24h	price_change_percentage_7d	price_change_perce
coin_id			
bitcoin	1.08388	7.60278	
ethereum	0.22392	10.38134	
tether	-0.21173	0.04935	
ripple	-0.37819	-0.60926	
bitcoin-cash	2.90585	17.09717	
binancecoin	2.10423	12.85511	
chainlink	-0.23935	20.69459	
cardano	0.00322	13.99302	
litecoin	-0.06341	6.60221	
bitcoin-cash-sv	0.92530	3.29641	

```
In [3]: 1 df_market_data.info()

<class 'pandas.core.frame.DataFrame'>
Index: 41 entries, bitcoin to digibyte
Data columns (total 7 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   price_change_percentage_24h           41 non-null     float64
1   price_change_percentage_7d            41 non-null     float64
2   price_change_percentage_14d           41 non-null     float64
3   price_change_percentage_30d           41 non-null     float64
4   price_change_percentage_60d           41 non-null     float64
5   price_change_percentage_200d          41 non-null     float64
6   price_change_percentage_1y            41 non-null     float64
dtypes: float64(7)
memory usage: 2.6+ KB
```


```
In [4]: 1 # Generate summary statistics
        2 df_market_data.describe()
```

Out[4]:

	price_change_percentage_24h	price_change_percentage_7d	price_change_percentage_1mo
count	41.000000	41.000000	41.000000
mean	-0.269686	4.497147	0.185147
std	2.694793	6.375218	8.376147
min	-13.527860	-6.094560	-18.158147
25%	-0.608970	0.047260	-5.026147
50%	-0.063410	3.296410	0.105147
75%	0.612090	7.602780	5.516147
max	4.840330	20.694590	24.238147

```
In [5]: 1 # we need a scaler due to variability in the max for each range
```

In [6]: ▶



Prepare the Data

```
In [7]: 1 # are we independent? or suffer from multi-collinearity?
2 corrs = df_market_data.corr()
3 corrs
```

```
Out[7]:
```

	price_change_percentage_24h	price_change_percentage_7d
price_change_percentage_24h	1.000000	0.169659
price_change_percentage_7d	0.169659	1.000000
price_change_percentage_14d	0.279682	0.538294
price_change_percentage_30d	0.292563	0.056899
price_change_percentage_60d	0.136974	-0.145099
price_change_percentage_200d	-0.541190	-0.052533
price_change_percentage_1y	-0.750630	-0.038424

```
In [8]: 1 sns.heatmap(corrs)
2 plt.show
```

```
Out[8]: <function matplotlib.pyplot.show(close=None, block=None)>
```



```
In [9]: 1 df_market_data.columns
```

```
Out[9]: Index(['price_change_percentage_24h', 'price_change_percentage_7d',
              'price_change_percentage_14d', 'price_change_percentage_30d',
              'price_change_percentage_60d', 'price_change_percentage_200d',
              'price_change_percentage_1y'],
              dtype='object')
```

```
In [10]: 1 # Use the `StandardScaler()` module from scikit-learn to normalize the
2 num_features = [ 'price_change_percentage_24h', 'price_change_percenta
3                 'price_change_percentage_14d', 'price_change_percentage_30d',
4                 'price_change_percentage_60d', 'price_change_percentage_200d',
5                 'price_change_percentage_1y']
```

In [11]:

```
1 # subset
2 df_sub = df_market_data.loc[:, num_features]
3
4 # initialize
5 scaler = StandardScaler()
6
7 # fit
8 scaler.fit(df_sub)
9
10 # predict/transform
11 scaled_data = scaler.transform(df_sub)
12 df_scaled = pd.DataFrame(scaled_data, columns=num_features)
13
14
15 df_scaled.head()
```

Out[11]:

	price_change_percentage_24h	price_change_percentage_7d	price_change_percentage_14d
0	0.508529	0.493193	0.772200
1	0.185446	0.934445	0.558692
2	0.021774	-0.706337	-0.021680
3	-0.040764	-0.810928	0.249458
4	1.193036	2.000959	1.760610

In [12]:

```
1 df_scaled.describe()
```

Out[12]:

	price_change_percentage_24h	price_change_percentage_7d	price_change_percentage_14d
count	41.000000	4.100000e+01	4.100000e+01
mean	0.000000	1.895503e-16	2.707861e-16
std	1.012423	1.012423e+00	1.012423e+00
min	-4.981042	-1.682027e+00	-2.217108e+00
25%	-0.127467	-7.066688e-01	-6.299628e-01
50%	0.077497	-1.906843e-01	-9.190922e-01
75%	0.331280	4.931931e-01	6.435649e-01
max	1.919812	2.572251e+00	2.907054e+00

Find the Best Value for k Using the Original Data.

```
In [13]: ▶ 1 # Create a list with the number of k-values from 1 to 11
          2 X = df_scaled.loc[:, num_features]
          3 X.head()
```

```
Out[13]:
```

	price_change_percentage_24h	price_change_percentage_7d	price_change_percentage_14d
0	0.508529	0.493193	0.772200
1	0.185446	0.934445	0.558692
2	0.021774	-0.706337	-0.021680
3	-0.040764	-0.810928	0.249458
4	1.193036	2.000959	1.760610


```
In [14]: ▶ 1 # Create a a List to store inertia values
          2 inertia = []
          3 silhouettes = []
          4 cha_chas = []
          5
          6 # Create a a List to store the values of k
          7 k = list(range(2, 11))
          8
          9 # Create a for-Loop where each value of k is evaluated using the K-means
         10 # Fit the model using the spread_df DataFrame
         11 # Append the value of the computed inertia from the `inertia_` attribute
         12 for i in k:
         13     # initialize the model
         14     k_model = KMeans(n_clusters=i, random_state=1)
         15
         16     # fit the model
         17     k_model.fit(X)
         18
         19     # predict the model
         20     preds = k_model.predict(X)
         21
         22     # evaluate the model (generate the metics)
         23     inertia.append(k_model.inertia_)
         24     score = silhouette_score(X, preds)
         25     silhouettes.append(score)
         26
         27     cha_cha = calinski_harabasz_score(X, preds)
         28     cha_chas.append(cha_cha)
         29
         30     print(f"Finished {i} out of {max(k)}")
```

```
In [15]: 1 # Define a DataFrame to hold the values for k and the corresponding i
2 elbow_data = {"k": k, "inertia": inertia, "silhouette_score": silhouette_score}
3 df_elbow = pd.DataFrame(elbow_data)
4
5 df_elbow["acc"] = df_elbow.inertia.diff()
6
7 # Review the DataFrame
8 df_elbow.head(10)
```

Out[15]:

	k	inertia	silhouette_score	cha_score	acc
0	2	195.820218	0.651576	18.159573	NaN
1	3	123.190482	0.702822	25.264783	-72.629736
2	4	79.022435	0.314482	32.459853	-44.168046
3	5	63.858668	0.329023	31.448698	-15.163768
4	6	53.057788	0.287883	30.864375	-10.800879
5	7	44.406791	0.290874	30.956861	-8.650998
6	8	37.078233	0.205692	31.776126	-7.328557
7	9	32.832187	0.258600	30.965687	-4.246046
8	10	28.165433	0.244422	31.653739	-4.666754


```
In [16]: ▶ 1 # Plot a line chart with all the inertia values computed with
2 # the different values of k to visually identify the optimal value for
3
4 # Plot the DataFrame
5 plt.plot(df_elbow["k"], df_elbow["inertia"])
6 plt.title("Elbow Curve")
7 plt.xticks(df_elbow["k"])
8 plt.ylabel("inertia")
9 plt.xlabel("k")
10 plt.show()
```



Answer the following question:

Question: What is the best value for k ?

Answer: 4 or 5

Cluster Cryptocurrencies with K-means Using the Original Data

```
In [17]: 1 # Initialize the K-Means model using the best value for k
        2 model = KMeans(n_clusters=4, random_state=1)
```

```
In [18]: 1 # Fit the K-Means model using the scaled data
        2 model.fit(X)
```

```
Out[18]: KMeans
KMeans(n_clusters=4, random_state=1)
```

```
In [19]: 1 # Predict the clusters to group the cryptocurrencies using the scaled
        2 preds = model.predict(X)
        3
```

```
In [20]: 1 # Add a new column to the DataFrame with the predicted clusters
        2 df2 = df_scaled.copy()
        3 df2['clusters'] = preds
        4
        5
        6 # Display sample data
        7 df2.head()
```

```
Out[20]:
```


	price_change_percentage_24h	price_change_percentage_7d	price_change_percentage_14d
0	0.508529	0.493193	0.772200
1	0.185446	0.934445	0.558692
2	0.021774	-0.706337	-0.021680

```

In [21]: ► 1 # Create a scatter plot using hvPlot by setting
2 # `x="price_change_percentage_24h"` and `y="price_change_percentage_7d"
3 # Color the graph points with the labels found using K-Means and
4 # add the crypto name in the `hover_cols` parameter to identify
5 # the cryptocurrency represented by each data point.
6 df2.hvplot.scatter(
7     width=800,
8     height=400,
9     x="price_change_percentage_24h",
10    y="price_change_percentage_7d",
11    by="clusters",
12    hover_cols="coin_id"
13 )

```

Out[21]:



```

In [22]: ► 1 df2.clusters.value_counts()

```

```

Out[22]: clusters
0      26
2      13
3       1
1       1
Name: count, dtype: int64

```

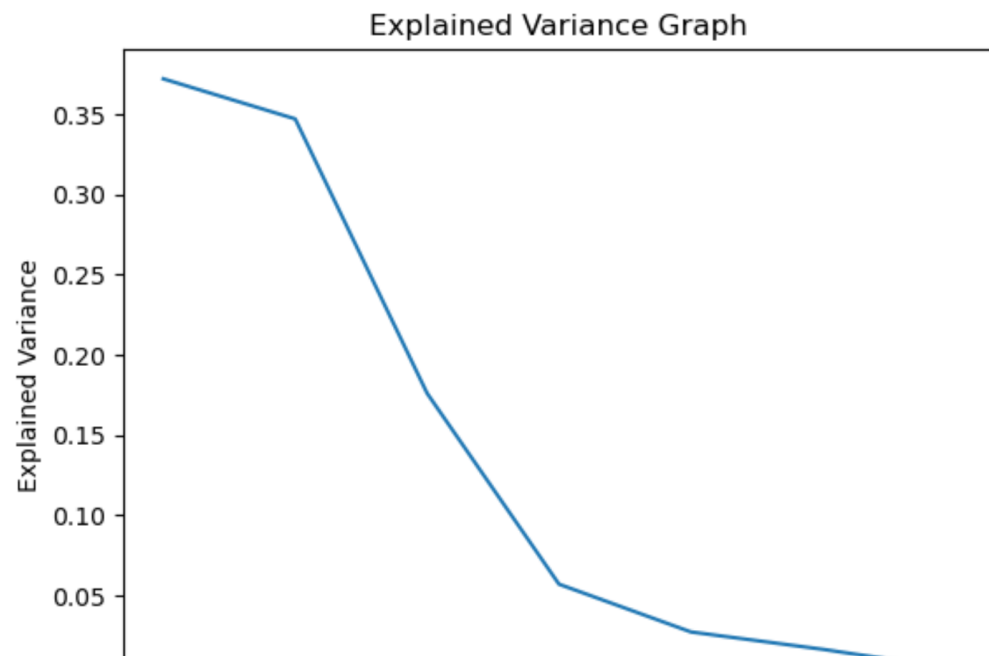
Optimize Clusters with Principal Component Analysis.

```
In [23]: 1 # Instantiate the PCA instance and declare the number of PCA variables
2 num_pca = len(num_features)
3 pca = PCA(n_components=num_pca)
4
5 # Fit the PCA model on the transformed credit card DataFrame
6 data_pca = pca.fit_transform(df_scaled.loc[:, num_features])
7
8 # Create the PCA DataFrame
9 df_pca = pd.DataFrame(
10     data_pca,
11     columns=[f"PCA{x+1}" for x in range(num_pca)]
12 )
13
14 df_pca.head()
```

Out[23]:

	PCA1	PCA2	PCA3	PCA4	PCA5	PCA6	PCA7
0	-0.600667	0.842760	0.461595	-0.109151	-0.033786	-0.225703	0.006595
1	-0.458261	0.458466	0.952877	0.095100	0.014588	0.034158	0.109593
2	-0.433070	-0.168126	-0.641752	-0.470282	0.115300	-0.127710	-0.086857
3	-0.471835	-0.222660	-0.479053	-0.737473	-0.148641	-0.273472	0.134870
4	-1.157800	2.041209	1.859715	0.236479	-0.191787	-0.411513	-0.070411

```
In [24]: ▶ 1 # Can we reduce the dimensions?
2
3 # Calculate the PCA explained variance ratio
4 exp_var = pca.explained_variance_ratio_
5
6 plt.plot(range(1, num_pca + 1), exp_var)
7 plt.title("Explained Variance Graph")
8 plt.xlabel("PCA #")
9 plt.ylabel("Explained Variance")
10 plt.xticks(range(1, num_pca + 1))
11 plt.show()
```



```
In [25]: ► 1 # Use the PCA model with `fit_transform` to reduce to
           2 # three principal components.
           3
           4 # View the first five rows of the DataFrame. (BOOOOTH Like 4 better)
           5 df3 = df_pca.loc[:, ["PCA1", "PCA2", "PCA3", "PCA4"]]
           6 df3.head()
```

Out[25]:

	PCA1	PCA2	PCA3	PCA4
0	-0.600667	0.842760	0.461595	-0.109151
1	-0.458261	0.458466	0.952877	0.095100
2	-0.433070	-0.168126	-0.641752	-0.470282
3	-0.471835	-0.222660	-0.479053	-0.737473
4	-1.157800	2.041209	1.859715	0.236479

```
In [26]: ► 1 print("Explained Variance")
           2 for i in range(len(exp_var)):
           3     val = exp_var[i]
           4     print(f"PCA{i+1}:", round(val, 3))
           5
           6 print()
           7 print("CUMULATIVE Explained Variance")
           8
           9 exp_var_cum = np.cumsum(exp_var)
          10 for i in range(len(exp_var_cum)):
          11     val = exp_var_cum[i]
          12     print(f"PCA{i+1}:", round(val, 3))
```

```
Explained Variance
PCA1: 0.372
PCA2: 0.347
PCA3: 0.176
PCA4: 0.057
```

```
In [27]: 1 pca.explained_variance_ratio_

Out[27]: array([0.3719856 , 0.34700813, 0.17603793, 0.05705673, 0.02729754,
               0.0164632 , 0.00415086])
```

```
In [28]: 1 pca.explained_variance_ratio_.sum()

Out[28]: 1.0
```

```
In [29]: 1 # Retrieve the explained variance to determine how much information
2 # can be attributed to each principal component.
3
```

Answer the following question:

Question: What is the total explained variance of the three principal components?

Answer: 0.895

Find the Best Value for k Using the PCA Data

```
In [31]: 1 # Create a list with the number of k-values from 1 to 11
2 X = df3.loc[:, ["PCA1", "PCA2", "PCA3", "PCA4"]]
3 X.head()
```

Out[31]:

	PCA1	PCA2	PCA3	PCA4
0	-0.600667	0.842760	0.461595	-0.109151
1	-0.458261	0.458466	0.952877	0.095100
2	-0.433070	-0.168126	-0.641752	-0.470282
3	-0.471835	-0.222660	-0.479053	-0.737473
4	-1.157800	2.041209	1.859715	0.236479

```
In [32]: ▶ 1 # Create a a list to store inertia values
2 inertia = []
3 silhouettes = []
4 cha_chas = []
5
6 # Create a a list to store the values of k
7 k = list(range(2, 11))
8
9 # Create a for-loop where each value of k is evaluated using the K-means
10 # Fit the model using the spread_df DataFrame
11 # Append the value of the computed inertia from the `inertia_` attribute
12 for i in k:
13     # initialize the model
14     k_model = KMeans(n_clusters=i, random_state=1)
15
16     # fit the model
17     k_model.fit(X)
18
19     # predict the model
20     preds = k_model.predict(X)
21
22     # evaluate the model (generate the metics)
23     inertia.append(k_model.inertia_)
24     score = silhouette_score(X, preds)
25     silhouettes.append(score)
26
27     cha_cha = calinski_harabasz_score(X, preds)
28     cha_chas.append(cha_cha)
29
30     print(f"Finished {i} out of {max(k)}")
```

```
Finished 2 out of 10
Finished 3 out of 10
Finished 4 out of 10
Finished 5 out of 10
```




```
In [33]: 1 # Define a DataFrame to hold the values for k and the corresponding ir
2 elbow_data = {"k": k, "inertia": inertia, "silhouette_score": silhouet
3 df_elbow2 = pd.DataFrame(elbow_data)
4
5 df_elbow["acc"] = df_elbow.inertia.diff()
6
7 # Review the DataFrame
8 df_elbow.head(10)
```

Out[33]:

	k	inertia	silhouette_score	cha_score	acc
0	2	195.820218	0.651576	18.159573	NaN
1	3	123.190482	0.702822	25.264783	-72.629736
2	4	79.022435	0.314482	32.459853	-44.168046
3	5	63.858668	0.329023	31.448698	-15.163768
4	6	53.057788	0.287883	30.864375	-10.800879
5	7	44.406791	0.290874	30.956861	-8.650998
6	8	37.078233	0.205692	31.776126	-7.328557
7	9	32.832187	0.258600	30.965687	-4.246046
8	10	28.165433	0.244422	31.653739	-4.666754

```
In [34]: ► 1 # Plot a line chart with all the inertia values computed with
2 # the different values of k to visually identify the optimal value for
3
4 # Plot the DataFrame
5 plt.plot(df_elbow2["k"], df_elbow["inertia"])
6 plt.title("Elbow Curve")
7 plt.xticks(df_elbow2["k"])
8 plt.ylabel("inertia")
9 plt.xlabel("k")
10 plt.show()
```



Answer the following questions:

- **Question:** What is the best value for `k` when using the PCA data?
 - **Answer:** 4
- **Question:** Does it differ from the best `k` value found using the original data?
 - **Answer:** No really, Just a bit clearer

Cluster Cryptocurrencies with K-means Using the PCA Data

```
In [35]: 1 # Initialize the K-Means model using the best value for k
        2 model = KMeans(n_clusters=2, random_state=1)
```

```
In [36]: 1 # Fit the K-Means model using the PCA data
        2 model.fit(X)
```

```
Out[36]: KMeans
KMeans(n_clusters=2, random_state=1)
```

```
In [37]: 1 # Predict the clusters to group the cryptocurrencies using the PCA data
        2 k_lower = model.predict(X)
        3 # Print the resulting array of cluster values.
        4
```

In [38]: ►


```
1 # Create a copy of the DataFrame with the PCA data
2 # Define the model with the higher value of k clusters
3 # Use a random_state of 1 to generate the model
4
5 # CHANGE THIS DEPENDING ON YOUR OPTIMAL k
6 model = KMeans(n_clusters=4, random_state=1)
7
8 # Fit the model
9 model.fit(X)
10
11 # Make predictions
12 preds = model.predict(X)
13
14 # Add a class column with the labels to the df DataFrame
15 df4 = df3.copy()
16 df4['clusters'] = preds
17
18 df4.head()
19
20 # Add a new column to the DataFrame with the predicted clusters
21
22
23 # Display sample data
24
```

Out[38]:

	PCA1	PCA2	PCA3	PCA4	clusters
0	-0.600667	0.842760	0.461595	-0.109151	2
1	-0.458261	0.458466	0.952877	0.095100	2
2	-0.433070	-0.168126	-0.641752	-0.470282	0
3	-0.471835	-0.222660	-0.479053	-0.737473	0

```
In [39]: ► 1 # Create a scatter plot using hvPlot by setting
2 # `x="PC1"` and `y="PC2"`.
3 # Color the graph points with the labels found using K-Means and
4 # add the crypto name in the `hover_cols` parameter to identify
5 # the cryptocurrency represented by each data point.
6 df4.hvplot.scatter(
7     width=800,
8     height=400,
9     x="PCA1",
10    y="PCA2",
11    by="clusters",
12    hover_cols="coin_id"
13 )
```

Out[39]:



Visualize and Compare the Results

In this section, you will visually analyze the cluster analysis results by contrasting the outcome with and without using the optimization techniques.

```
In [40]: ▶ 1 # Composite plot to contrast the Elbow curves
2 # Create the Elbow curve plots for the original and second Elbow curve
3 elbow_plot1 = df_elbow.hvplot.line(x='k', y='inertia', title='Original')
4 elbow_plot2 = df_elbow2.hvplot.line(x='k', y='inertia', title='Second')
5
6 # Create a composite plot to contrast the Elbow curves
7 composite_elbow_plot = elbow_plot1 + elbow_plot2
8
9 # Display the composite plot
10 composite_elbow_plot
```

Out[40]:



Answer the following question:

- **Question:** After visually analyzing the cluster analysis results, what is the impact of using fewer features to cluster the data using K-Means?
- ****Answer:**** DATA IS MORE CONSCISE AND LESS DISPERSE