

## Preprocessing

```
In [1]: # Import the required modules
1 import pandas as pd
2 pd.set_option('display.max_columns', None)
3 import numpy as np
4
5 # visualization
6 import matplotlib.pyplot as plt
7 import seaborn as sns
8
9 # Machine Learning
10 from sklearn.model_selection import train_test_split
11
12 # Algorithms
13 from sklearn.linear_model import LogisticRegression
14 from sklearn.ensemble import RandomForestClassifier
15 from lightgbm import LGBMClassifier
16 from xgboost import XGBClassifier
17
18 # tensorflow
19 import tensorflow as tf
20
21 # Metrics
22 from sklearn.metrics import confusion_matrix, classification_report
23 from sklearn.metrics import roc_curve, roc_auc_score
24
25 # Preprocessing
26 from sklearn.preprocessing import StandardScaler
27 from sklearn.decomposition import PCA
28
29 # suppress warnings
30 import warnings
31 warnings.filterwarnings('ignore')
```

```
In [2]: # Import and read the charity_data.csv.  
2 df = pd.read_csv("https://static.bc-edx.com/data/dl-1-2/m21/lms/starter/  
3 df.head()
```

Out[2]:

	EIN	NAME	APPLICATION_TYPE	AFFILIATION	CLASSIFICATION	USE_
0	10520599	BLUE KNIGHTS MOTORCYCLE CLUB	T10	Independent	C1000	Prod
1	10531628	AMERICAN CHESAPEAKE CLUB CHARITABLE TR	T3	Independent	C2000	Prese
2	10547893	ST CLOUD PROFESSIONAL FIREFIGHTERS	T5	CompanySponsored	C3000	Prod
3	10553066	SOUTHSIDE ATHLETIC ASSOCIATION	T3	CompanySponsored	C2000	Prese
4	10556103	GENETIC RESEARCH INSTITUTE OF THE DESERT	T3	Independent	C1000	He

```
In [3]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>  
RangeIndex: 34299 entries, 0 to 34298  
Data columns (total 12 columns):  
 #   Column           Non-Null Count  Dtype    
---  --    
 0   EIN             34299 non-null   int64   
 1   NAME            34299 non-null   object   
 2   APPLICATION_TYPE 34299 non-null   object   
 3   AFFILIATION     34299 non-null   object   
 4   CLASSIFICATION  34299 non-null   object   
 5   USE_            34299 non-null   object
```

```
In [4]: 1 df.describe()
```

Out[4]:

	EIN	STATUS	ASK_AMT	IS_SUCCESSFUL
count	3.429900e+04	34299.000000	3.429900e+04	34299.000000
mean	5.191852e+08	0.999854	2.769199e+06	0.532406
std	2.451472e+08	0.012073	8.713045e+07	0.498956
min	1.052060e+07	0.000000	5.000000e+03	0.000000
25%	2.748482e+08	1.000000	5.000000e+03	0.000000
50%	4.656317e+08	1.000000	5.000000e+03	1.000000
75%	7.526117e+08	1.000000	7.742000e+03	1.000000
max	9.960869e+08	1.000000	8.597806e+09	1.000000

```
In [5]: 1 df.IS_SUCCESSFUL.value_counts()
```

Out[5]: IS\_SUCCESSFUL

```
1    18261  
0    16038  
Name: count, dtype: int64
```

```
In [6]: 1 # Drop the non-beneficial ID columns, 'EIN' and 'NAME'.  
2 df2 = df.copy()  
3  
4 df2 = df2.drop(["EIN", "NAME"], axis=1)  
5 df2.head()
```

Out[6]:

	APPLICATION_TYPE	AFFILIATION	CLASSIFICATION	USE_CASE	ORGANIZATION	ST/
0	T10	Independent	C1000	ProductDev	Association	
1	T3	Independent	C2000	Preservation	Co-operative	

```
In [7]: # Determine the number of unique values in each column.  
Out[7]: df2.nunique()
```

```
APPLICATION_TYPE      17  
AFFILIATION          6  
CLASSIFICATION      71  
USE_CASE              5  
ORGANIZATION         4  
STATUS                2  
INCOME_AMT           9  
SPECIAL_CONSIDERATIONS  2  
ASK_AMT             8747  
IS_SUCCESSFUL        2  
dtype: int64
```

```
In [8]: # Look at APPLICATION_TYPE value counts for binning  
Out[8]: application_type_counts = df2['APPLICATION_TYPE'].value_counts()  
print(application_type_counts)
```

```
APPLICATION_TYPE  
T3      27037  
T4      1542  
T6      1216  
T5      1173  
T19     1065  
T8      737  
T7      725  
T10     528  
T9      156  
T13     66  
T12     27  
T2      16  
T25     3  
T14     3  
T29     2  
T15     2  
T17     1  
Name: count, dtype: int64
```

```
In [9]: 1 df2['APPLICATION_TYPE'].value_counts() < 500
```

```
Out[9]: APPLICATION_TYPE
T3      False
T4      False
T6      False
T5      False
T19     False
T8      False
T7      False
T10     False
T9      True
T13     True
T12     True
T2      True
T25     True
T14     True
T29     True
T15     True
T17     True
Name: count, dtype: bool
```

```
In [10]: 1 # Choose a cutoff value and create a list of application types to be rep
2 # use the variable name `application_types_to_replace`
3 application_types_to_replace = ["T9", "T13", "T12", "T2", "T25", "T14",
4
5 # Replace in dataframe
6 for app in application_types_to_replace:
7     df2['APPLICATION_TYPE'] = df2['APPLICATION_TYPE'].replace(app,"Other")
8
9 # Check to make sure binning was successful
10 df2['APPLICATION_TYPE'].value_counts()
```

```
Out[10]: APPLICATION_TYPE
T3      27037
T4      1542
T6      1216
```

```
In [11]: # Look at CLASSIFICATION value counts for binning
          classification_counts = df2['CLASSIFICATION'].value_counts()
          print(classification_counts)
```

```
CLASSIFICATION
C1000      17326
C2000      6074
C1200      4837
C3000      1918
C2100      1883
...
C4120       1
C8210       1
C2561       1
C4500       1
C2150       1
Name: count, Length: 71, dtype: int64
```

```
In [13]: # You may find it helpful to look at CLASSIFICATION value counts >1
          classification_counts = df2['CLASSIFICATION'].value_counts()
          classification_counts_gt1 = classification_counts[classification_counts
          5]
```

```
CLASSIFICATION
C1000      17326
C2000      6074
C1200      4837
C3000      1918
C2100      1883
C7000       777
C1700       287
C4000       194
C5000       116
C1270       114
C2700       104
```

```
In [12]: # Choose a cutoff value and create a List of classifications to be replaced
# use the variable name `classifications_to_replace`
classifications_to_replace = list(df2['CLASSIFICATION'].value_counts().index)
# Replace in dataframe
for cls in classifications_to_replace:
    df2['CLASSIFICATION'] = df2['CLASSIFICATION'].replace(cls,"Other")
# Check to make sure binning was successful
df2['CLASSIFICATION'].value_counts()
```

```
Out[12]: CLASSIFICATION
C1000      17326
C2000       6074
C1200       4837
C3000       1918
C2100       1883
Other        1484
C7000        777
Name: count, dtype: int64
```

```
In [13]: # Convert categorical data to numeric with `pd.get_dummies`
df2 = pd.get_dummies(df2)
df2.head()
```

```
Out[13]:
```

	STATUS	ASK_AMT	IS_SUCCESSFUL	APPLICATION_TYPE_Other	APPLICATION_TYPE_T10
0	1	5000	1	False	True
1	1	108590	1	False	False
2	1	5000	0	False	False
3	1	6692	1	False	False

```
In [14]: # Split our preprocessed data into our features and target arrays
1 X = df2.drop('IS_SUCCESSFUL', axis=1)
2 y = df2['IS_SUCCESSFUL']
3
4 # Split the preprocessed data into a training and testing dataset
5 X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2,
6
7 print(X_train.shape)
8 print(X_test.shape)
```

(27439, 44)  
(6860, 44)

```
In [15]: # Create a StandardScaler instances
1 scaler = StandardScaler()
2
3 # Fit the StandardScaler
4 X_scaler = scaler.fit(X_train)
5
6 # Scale the data
7 X_train_scaled = X_scaler.transform(X_train)
8 X_test_scaled = X_scaler.transform(X_test)
```

## Compile, Train and Evaluate the Model

```
In [17]: # Define the model - deep neural net, i.e., the number of input features  
# YOUR CODE GOES HERE  
  
3  
4 nn1 = tf.keras.models.Sequential()  
5  
6 # First hidden Layer  
7 nn1.add(tf.keras.layers.Dense(units=5, activation='relu', input_dim=len(  
8  
9 # Second hidden Layer  
10 nn1.add(tf.keras.layers.Dense(units=3, activation='relu'))  
11  
12 # Output Layer  
13 nn1.add(tf.keras.layers.Dense(units=1, activation='sigmoid'))  
14  
15  
16 # Check the structure of the model  
17 nn1.summary()
```

Model: "sequential\_1"

Layer (type)	Output Shape
dense_3 (Dense)	(None, 5)
dense_4 (Dense)	(None, 3)
dense_5 (Dense)	(None, 1)

Total params: 247 (988.00 B)

```
Total params: 247 (988.00 B)
```

```
Trainable params: 247 (988.00 B)
```

```
Non-trainable params: 0 (0.00 B)
```

```
In [18]: # Compile the model  
nn.compile(loss='binary_crossentropy', optimizer='adam', metrics=['accuracy'])
```

```
In [20]: # Train the model  
fit_model = nn.fit(X_train_scaled, y_train, epochs=30)  
  
Epoch 1/30  
858/858 1s 1ms/step - accuracy: 0.7230 - loss: 0.5549  
Epoch 2/30  
858/858 1s 1ms/step - accuracy: 0.7205 - loss: 0.5602  
Epoch 3/30  
858/858 1s 1ms/step - accuracy: 0.7290 - loss: 0.5550  
Epoch 4/30  
858/858 1s 1ms/step - accuracy: 0.7255 - loss: 0.5553  
Epoch 5/30  
858/858 1s 1ms/step - accuracy: 0.7239 - loss: 0.5558  
Epoch 6/30  
858/858 1s 1ms/step - accuracy: 0.7232 - loss: 0.5587  
Epoch 7/30  
858/858 1s 1ms/step - accuracy: 0.7299 - loss: 0.5523  
Epoch 8/30  
858/858 1s 1ms/step - accuracy: 0.7305 - loss: 0.5500  
Epoch 9/30  
858/858 1s 1ms/step - accuracy: 0.7265 - loss: 0.5556  
Epoch 10/30  
858/858 1s 1ms/step - accuracy: 0.7252 - loss: 0.5539  
Epoch 11/30  
858/858 1s 1ms/step - accuracy: 0.7361 - loss: 0.5452  
Epoch 12/30  
858/858 1s 1ms/step - accuracy: 0.7312 - loss: 0.5513  
Epoch 13/30  
858/858 1s 1ms/step - accuracy: 0.7277 - loss: 0.5534  
Epoch 14/30  
858/858 1s 2ms/step - accuracy: 0.7218 - loss: 0.5582
```

```
In [21]: # Evaluate the model using the test data
          2 model_loss, model_accuracy = nn.evaluate(X_test_scaled,y_test,verbose=2)
          3 print(f"Loss: {model_loss}, Accuracy: {model_accuracy}")


```

```
215/215 - 0s - 1ms/step - accuracy: 0.7198 - loss: 0.5620
Loss: 0.562018871307373, Accuracy: 0.7198250889778137
```

```
In [22]: def doClassification_NN(nn_model, X_train, X_test, y_train, y_test):
          2     # predict
          3     train_preds = nn_model.predict(X_train)
          4     train_preds = tf.greater(train_preds, .5)
          5     train_probs = nn_model.predict(X_train)
          6
          7     test_preds = nn_model.predict(X_test)
          8     test_preds = tf.greater(test_preds, .5)
          9     test_probs = nn_model.predict(X_test)
         10
         11     # evaluate train
         12     train_cr = classification_report(y_train, train_preds)
         13     train_cm = confusion_matrix(y_train, train_preds)
         14
         15     train_report = f"""
         16     Train Confusion Matrix:
         17     {train_cm}
         18
         19     Train Report:
         20     {train_cr}
         21     """
         22
         23     print("TRAINING METRICS")
         24     print(train_report)
         25     print()
         26
         27     # train ROC curve
         28     # Compute fpr, tpr, thresholds and roc auc
         29     fpr, tpr, thresholds = roc_curve(y_train, train_probs)
         30     roc_auc = roc_auc_score(y_train, train_probs)
```

```

30
31     # Plot ROC curve
32     plt.plot(fpr, tpr, label='ROC curve (area = %0.3f)' % roc_auc)
33     plt.plot([0, 1], [0, 1], 'k--') # random predictions curve
34     plt.xlim([0.0, 1.0])
35     plt.ylim([0.0, 1.0])
36     plt.xlabel('False Positive Rate or (1 - Specificity)')
37     plt.ylabel('True Positive Rate or (Sensitivity)')
38     plt.title('TRAINING Receiver Operating Characteristic')
39     plt.legend(loc="lower right")
40     plt.show()
41     print()
42     print()
43
44     # evaluate test
45     test_cr = classification_report(y_test, test_preds)
46     test_cm = confusion_matrix(y_test, test_preds)
47
48     test_report = f"""
49     Test Confusion Matrix:
50     {test_cm}
51
52     Test Report:
53     {test_cr}
54     """
55     print("TESTING METRICS")
56     print(test_report)
57     print()
58
59     # train ROC curve
60     # Compute fpr, tpr, thresholds and roc auc
61     fpr, tpr, thresholds = roc_curve(y_test, test_probs)
62     roc_auc = roc_auc_score(y_test, test_probs)
63
64     # Plot ROC curve
65     plt.plot(fpr, tpr, label='ROC curve (area = %0.3f)' % roc_auc)
66     plt.plot([0, 1], [0, 1], 'k--') # random predictions curve
67     plt.xlim([0.0, 1.0])
68     plt.ylim([0.0, 1.0])
69     plt.xlabel('False Positive Rate or (1 - Specificity)')

```

```
71 |     plt.title('TESTING Receiver Operating Characteristic')
72 |     plt.legend(loc="lower right")
73 |     plt.show()
```

In [32]: 1 doClassification\_NN(nn, X\_train\_scaled, X\_test\_scaled, y\_train, y\_test)

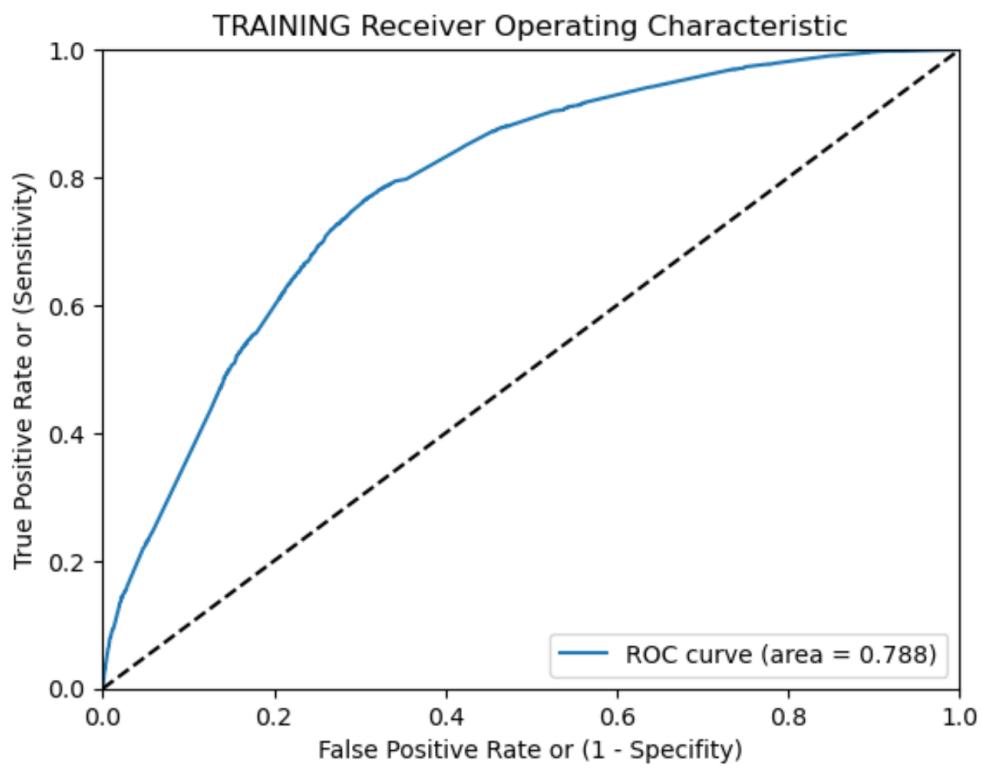
```
858/858 ━━━━━━ 1s 1ms/step
858/858 ━━━━━━ 1s 1ms/step
215/215 ━━━━ 0s 1ms/step
215/215 ━━━━ 0s 1ms/step
TRAINING METRICS
```

Train Confusion Matrix:

```
[[ 8695  4147]
 [ 3199 11398]]
```

Train Report:

	precision	recall	f1-score	support
0	0.73	0.68	0.70	12842
1	0.73	0.78	0.76	14597
accuracy			0.73	27439
macro avg	0.73	0.73	0.73	27439
weighted avg	0.73	0.73	0.73	27439

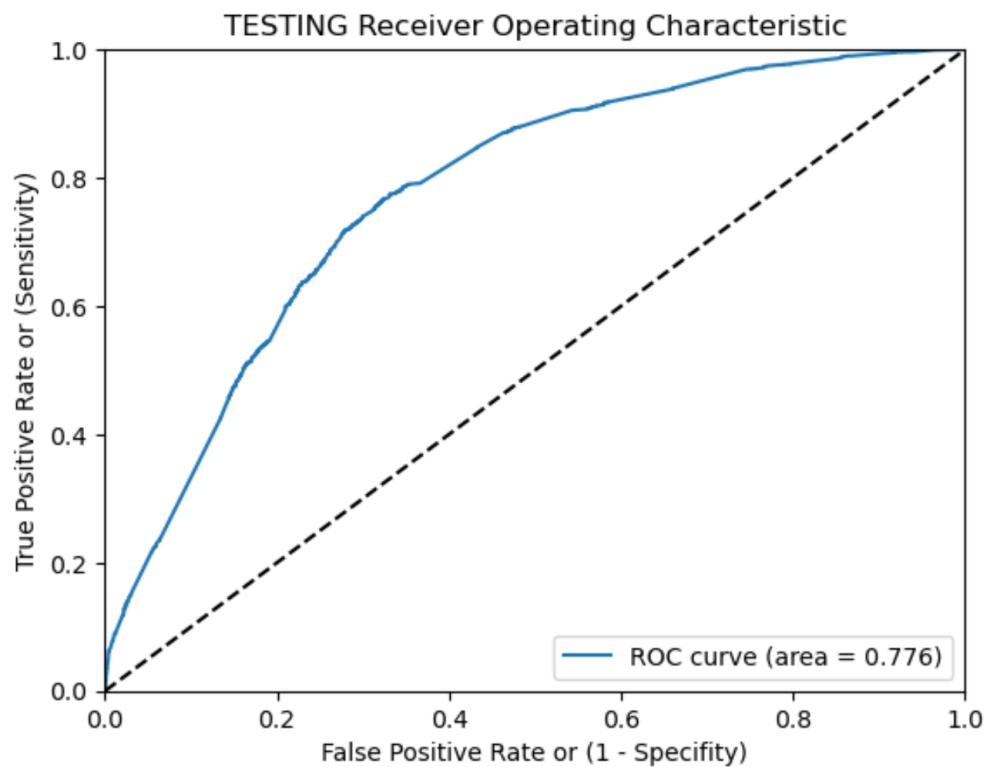


### TESTING METRICS

Test Confusion Matrix:  
[[2137 1059]  
 [ 826 2838]]

Test Report:

	precision	recall	f1-score	support
0	0.72	0.67	0.69	3196
1	0.73	0.77	0.75	3664
accuracy			0.73	6860
macro avg	0.72	0.72	0.72	6860
weighted avg	0.72	0.73	0.72	6860



```
In [23]: 1 # Export our model to HDF5 file
          2 nn1.save("nn1.h5")

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

In [24]: 1 # OPTIMIZATIONS

In [25]: 1 # Define the model - deep neural net, i.e., the number of input features
          2
          3 nn2 = tf.keras.models.Sequential()
          4
          5 # Add our first Dense Layer, including the input layer
          6 nn2.add(tf.keras.layers.Dense(units=15, activation="relu", input_dim=len)
          7
          8 # Add a second layer
          9 nn2.add(tf.keras.layers.Dense(units=7, activation="relu"))
          10
          11 # Add a third layer
          12 nn2.add(tf.keras.layers.Dense(units=5, activation="relu"))
          13
          14 # Add the output layer that uses a probability activation function
          15 nn2.add(tf.keras.layers.Dense(units=1, activation="sigmoid"))
          16
          17 # Check the structure of the Sequential model
          18 nn2.summary()
```

Model: "sequential\_2"

Layer (type)	Output Shape
dense_6 (Dense)	(None, 15)
dense_7 (Dense)	(None, 7)
dense_8 (Dense)	(None, 5)
dense_9 (Dense)	(None, 1)

Total params: 833 (3.25 KB)

Trainable params: 833 (3.25 KB)

Non-trainable params: 0 (0.00 B)

In [27]:

```
1 # Compile the Sequential model together and customize metrics
2 nn2.compile(loss="binary_crossentropy", optimizer="adam", metrics=["accu
3
4 # Fit the model to the training data
5 fit_model = nn2.fit(X_train_scaled, y_train, epochs=30)
```

```
Epoch 1/30
858/858 ━━━━━━━━ 2s 1ms/step - accuracy: 0.6624 - loss: 0.6312
Epoch 2/30
858/858 ━━━━━━ 1s 1ms/step - accuracy: 0.7301 - loss: 0.5561
Epoch 3/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7328 - loss: 0.5515
Epoch 4/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7304 - loss: 0.5532
Epoch 5/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7284 - loss: 0.5566
Epoch 6/30
858/858 ━━━━ 1s 1ms/step - accuracy: 0.7284 - loss: 0.5532
Epoch 7/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7321 - loss: 0.5467
Epoch 8/30
858/858 ━━━━ 1s 1ms/step - accuracy: 0.7368 - loss: 0.5417
Epoch 9/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7314 - loss: 0.5464
Epoch 10/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7325 - loss: 0.5475
Epoch 11/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7301 - loss: 0.5510
Epoch 12/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7334 - loss: 0.5472
Epoch 13/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7333 - loss: 0.5485
Epoch 14/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7337 - loss: 0.5440
Epoch 15/30
858/858 ━━━━ 1s 2ms/step - accuracy: 0.7347 - loss: 0.5466
Epoch 16/30
858/858 ━━━━ 1s 1ms/step - accuracy: 0.7332 - loss: 0.5462
```

```
In [28]: # Evaluate the model using the test data
1 model_loss, model_accuracy = nn2.evaluate(X_test_scaled, y_test,verbose=1)
2
3 print(f"Loss: {model_loss}, Accuracy: {model_accuracy}")
```

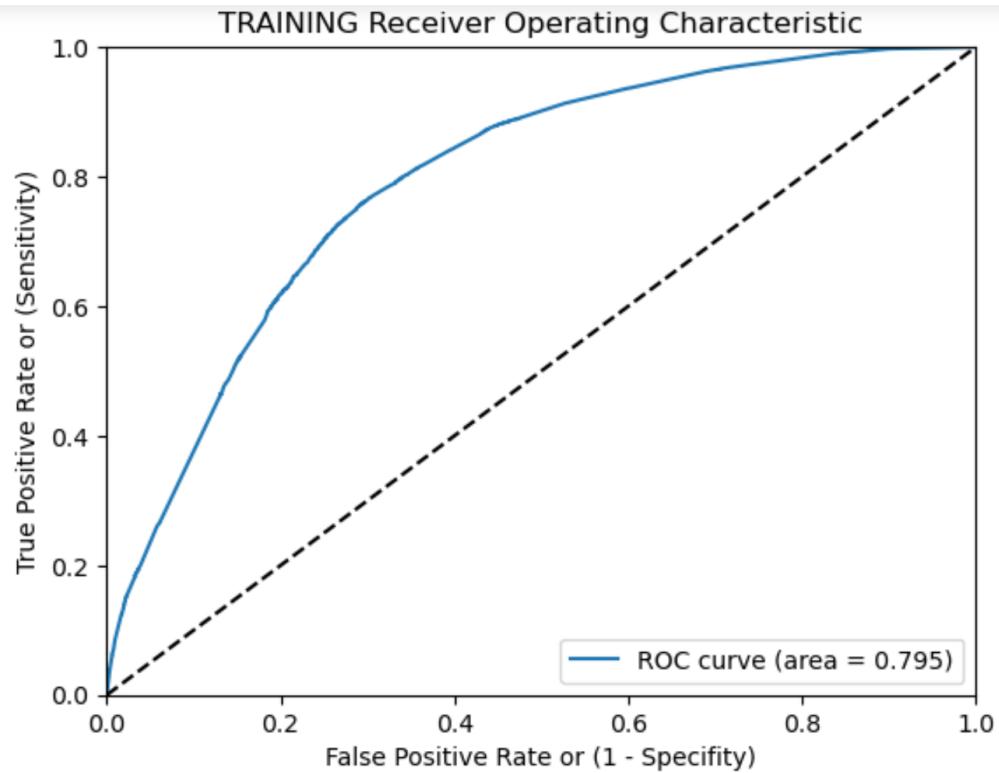
```
215/215 - 0s - 2ms/step - accuracy: 0.7216 - loss: 0.5547
Loss: 0.5547159910202026, Accuracy: 0.7215743660926819
```

```
In [30]: doClassification_NN(nn2, X_train_scaled, X_test_scaled, y_train, y_test)
```

```
858/858 ━━━━━━━━ 1s 1ms/step
858/858 ━━━━━━━━ 1s 1ms/step
215/215 ━━━━━━ 0s 1ms/step
215/215 ━━━━━━ 0s 1ms/step
TRAINING METRICS
```

```
Train Confusion Matrix:
[[ 8594 4248]
 [ 3077 11520]]
```

```
Train Report:
precision    recall    f1-score    support
          0       0.74      0.67      0.70     12842
          1       0.73      0.79      0.76     14597
accuracy                           0.73     27439
macro avg       0.73      0.73      0.73     27439
weighted avg    0.73      0.73      0.73     27439
```



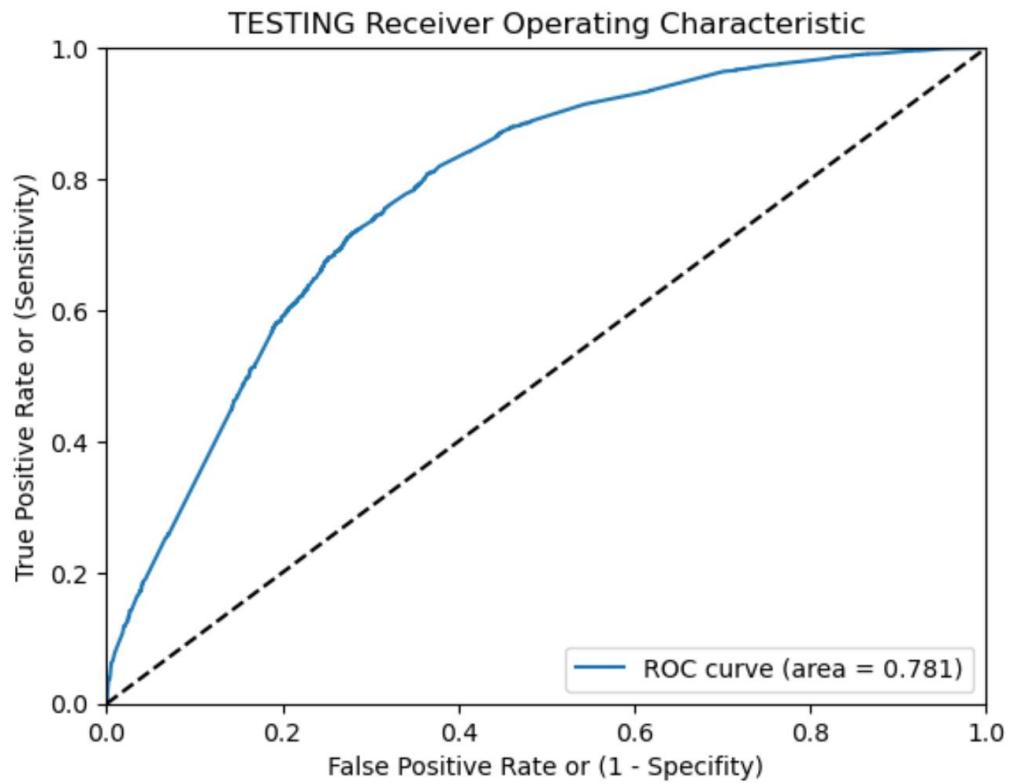
#### TESTING METRICS

Test Confusion Matrix:  
[[2080 1116]  
[ 794 2870]]

Test Report:

	precision	recall	f1-score	support
0	0.72	0.65	0.69	3196
1	0.72	0.78	0.75	3664

accuracy			0.72		6860
macro avg	0.72	0.72	0.72	6860	
weighted avg	0.72	0.72	0.72	6860	



```
In [31]: # Export our model to HDF5 file
2 nn2.save("nn2.h5")

WARNING:absl:You are saving your model as an HDF5 file via `model.save()` or `keras.saving.save_model(model)`. This file format is considered legacy. We recommend using instead the native Keras format, e.g. `model.save('my_model.keras')` or `keras.saving.save_model(model, 'my_model.keras')`.

In [34]: # Define the model - deep neural net, i.e., the number of input features
2
3 nn3 = tf.keras.models.Sequential()
4
5 # Add our first Dense layer, including the input layer
6 nn3.add(tf.keras.layers.Dense(units=25, activation="relu", input_dim=len
7
8 # Add a second layer
9 nn3.add(tf.keras.layers.Dense(units=14, activation="relu"))
10
11 # Add a third layer
12 nn3.add(tf.keras.layers.Dense(units=7, activation="relu"))
13
14 # Add the output layer that uses a probability activation function
15 nn3.add(tf.keras.layers.Dense(units=1, activation="sigmoid"))
16
17 # Check the structure of the Sequential model
18 nn3.summary()
```

Model: "sequential\_3"

Layer (type)	Output Shape
dense_10 (Dense)	(None, 25)
dense_11 (Dense)	(None, 14)
dense_12 (Dense)	(None, 7)
dense_13 (Dense)	(None, 1)

Total params: 1,602 (6.26 KB)

Trainable params: 1,602 (6.26 KB)

Non-trainable params: 0 (0.00 B)

```
In [32]: ❶ def doClassification(model, X_train, X_test, y_train, y_test):
    # predict
    train_preds = model.predict(X_train)
    train_probs = model.predict_proba(X_train)

    test_preds = model.predict(X_test)
    test_probs = model.predict_proba(X_test)

    # evaluate train
    train_cr = classification_report(y_train, train_preds)
    train_cm = confusion_matrix(y_train, train_preds)

    train_report = f"""
Train Confusion Matrix:
{train_cm}

Train Report:
{train_cr}
"""

    print("TRAINING METRICS")
    print(train_report)
    print()

    # train ROC curve
    # Compute fpr, tpr, thresholds and roc auc
    fpr, tpr, thresholds = roc_curve(y_train, train_probs[:,1])
    roc_auc = roc_auc_score(y_train, train_probs[:,1])

    # Plot ROC curve
    plt.plot(fpr, tpr, label='ROC curve (area = %0.3f)' % roc_auc)
    plt.plot([0, 1], [0, 1], 'k--') # random predictions curve
    plt.xlim([0.0, 1.0])
    plt.ylim([0.0, 1.0])
    plt.xlabel('False Positive Rate or (1 - Specificity)')
    plt.ylabel('True Positive Rate or (Sensitivity)')
    plt.title('TRAINING Receiver Operating Characteristic')
    plt.legend(loc="lower right")
    plt.show()
    print()
    print()
```

```
42 # evaluate test
43 test_cr = classification_report(y_test, test_preds)
44 test_cm = confusion_matrix(y_test, test_preds)
45
46 test_report = f"""
47 Test Confusion Matrix:
48 {test_cm}
49
50 Test Report:
51 {test_cr}
52 """
53 print("TESTING METRICS")
54 print(test_report)
55 print()
56
57 # train ROC curve
58 # Compute fpr, tpr, thresholds and roc auc
59 fpr, tpr, thresholds = roc_curve(y_test, test_probs[:,1])
60 roc_auc = roc_auc_score(y_test, test_probs[:,1])
61
62 # Plot ROC curve
63 plt.plot(fpr, tpr, label='ROC curve (area = %0.3f)' % roc_auc)
64 plt.plot([0, 1], [0, 1], 'k--') # random predictions curve
65 plt.xlim([0.0, 1.0])
66 plt.ylim([0.0, 1.0])
67 plt.xlabel('False Positive Rate or (1 - Specificity)')
68 plt.ylabel('True Positive Rate or (Sensitivity)')
69 plt.title('TESTING Receiver Operating Characteristic')
70 plt.legend(loc="lower right")
71 plt.show()
```

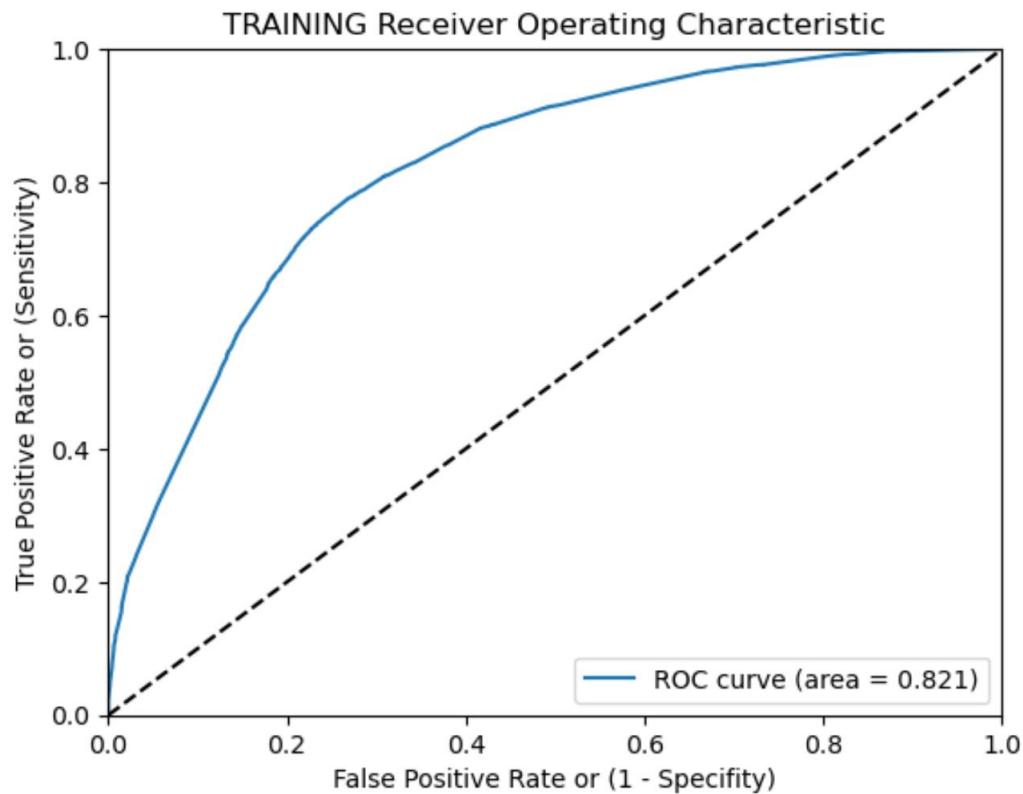
```
In [33]: 1 # initialize
          2 xgb = XGBClassifier(random_state=42)
          3
          4 # fit
          5 xgb.fit(X_train_scaled, y_train)
          6
          7 doClassification(xgb, X_train_scaled, X_test_scaled, y_train, y_test)
```

#### TRAINING METRICS

Train Confusion Matrix:  
[[ 9151 3691]  
[ 3062 11535]]

Train Report:

	precision	recall	f1-score	support
0	0.75	0.71	0.73	12842
1	0.76	0.79	0.77	14597
accuracy			0.75	27439
macro avg	0.75	0.75	0.75	27439
weighted avg	0.75	0.75	0.75	27439



## TESTING METRICS

Test Confusion Matrix:

```
[[2148 1048]
 [ 874 2790]]
```

Test Report:

	precision	recall	f1-score	support
0	0.71	0.67	0.69	3196
1	0.73	0.76	0.74	3664
accuracy			0.72	6860
macro avg	0.72	0.72	0.72	6860
weighted avg	0.72	0.72	0.72	6860

